

# Anomaly Detection

5/21

# Outline

- What is Anomaly Detection
- Classic Method
  - With Classifier
  - GMM (Gaussian Mixture Model)
  - Auto-Encoder
  - PCA
  - Isolation Forest
  - Summary
- Anomaly Detection on image
  - AnoGAN
  - EGBAD
  - GANomaly
  - Summary
- Anomaly Detection on Audio
  - GMGAN


# What is Anomaly Detection

# What is Anomaly

- Training Data



Anomaly



Anomaly

# Classic Method

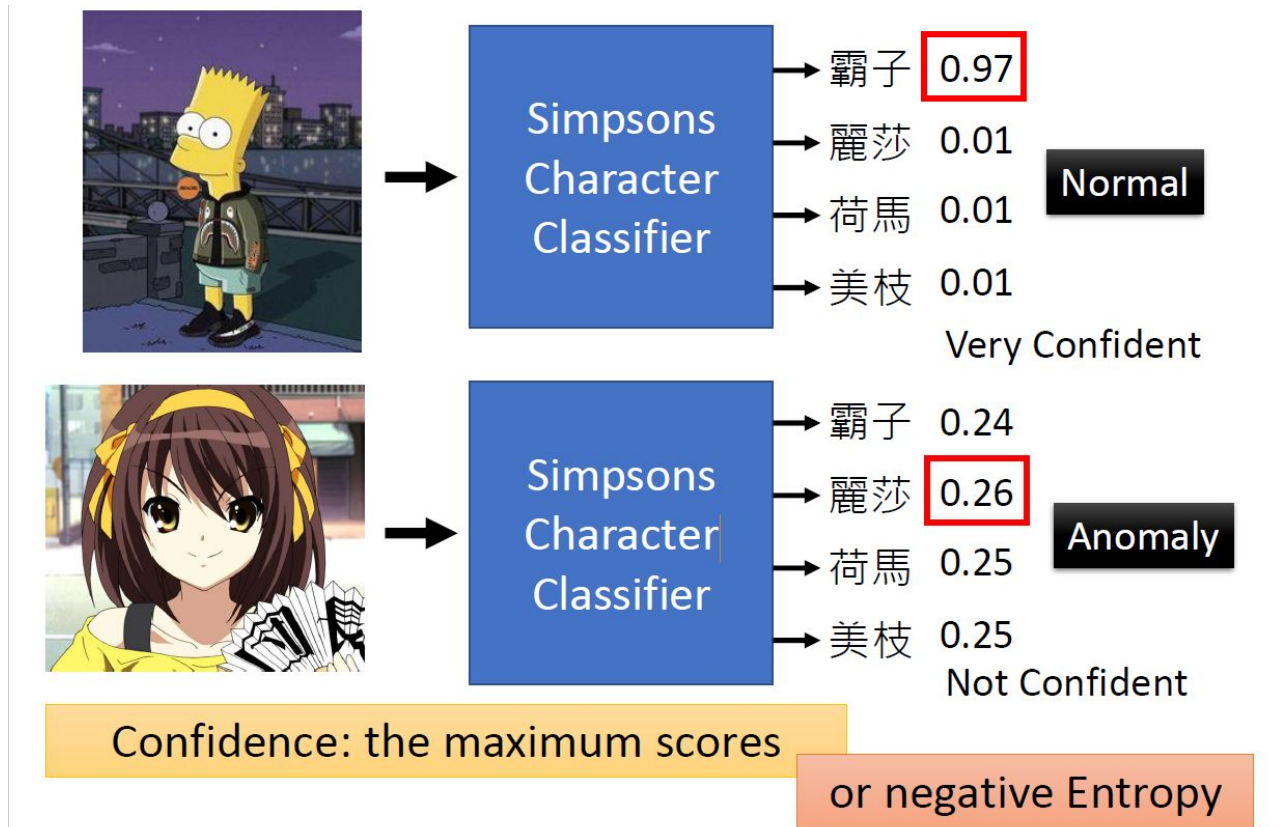
## With Classifier



### **Anomaly Detection:**


$$f(x) = \begin{cases} normal, & c(x) > \lambda \\ anomaly, & c(x) \leq \lambda \end{cases}$$

# With Classifier



# GMM (Gaussian Mixture Model)

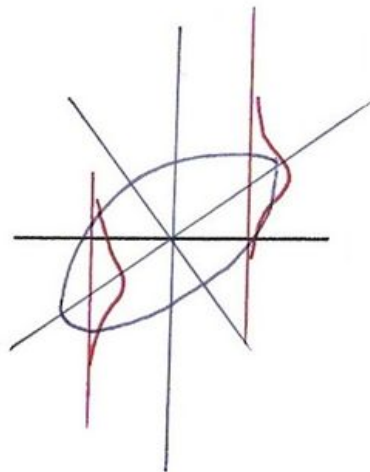
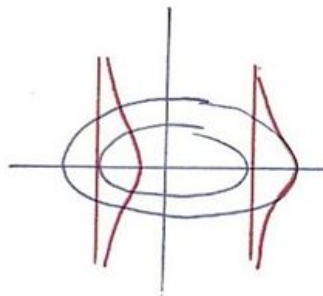
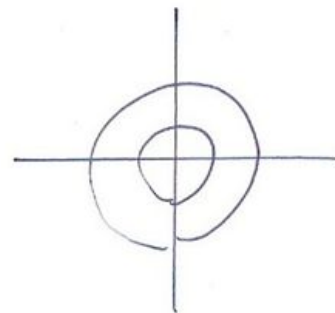

## 1-dim Gaussian Mixtures





# GMM (Gaussian Mixture Model)

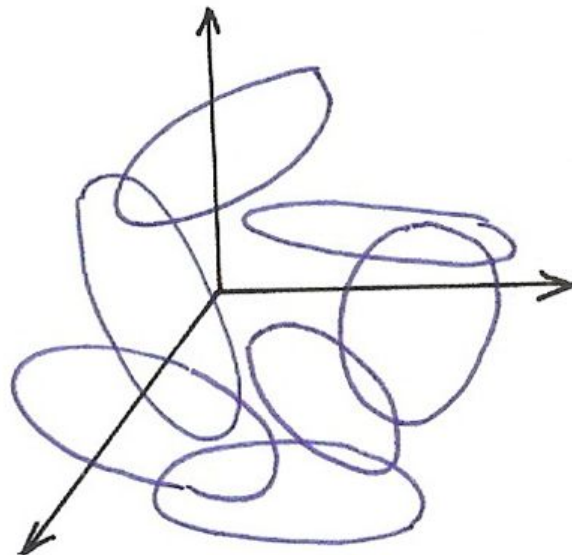
## 2-dim Gaussian



# GMM (Gaussian Mixture Model)

N-dim Gaussian Mixtures

N-dim Gaussian Mixtures

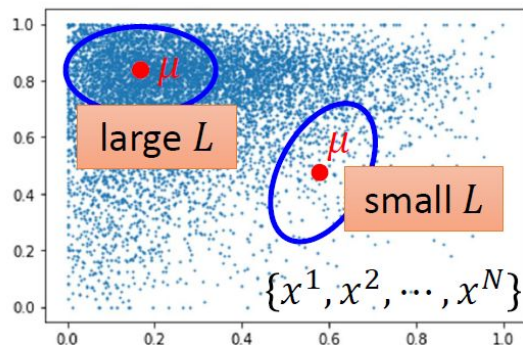


# GMM (Gaussian Mixture Model)

$$f_{\mu, \Sigma}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

Input: vector  $x$ , output: probability of sampling  $x$

$\theta$  which determines the shape of the function are **mean  $\mu$**   
and **covariance matrix  $\Sigma$**



$$L(\theta) = f_{\theta}(x^1) f_{\theta}(x^2) \cdots f_{\theta}(x^N)$$

$$L(\mu, \Sigma) = f_{\mu, \Sigma}(x^1) f_{\mu, \Sigma}(x^2) \cdots f_{\mu, \Sigma}(x^N)$$

$$\theta^* = \arg \max_{\theta} L(\theta)$$

$$\mu^*, \Sigma^* = \arg \max_{\mu, \Sigma} L(\mu, \Sigma)$$

$$\mu^* = \frac{1}{N} \sum_{n=1}^N x^n = \begin{bmatrix} 0.29 \\ 0.73 \end{bmatrix}$$

$$\Sigma^* = \frac{1}{N} \sum_{n=1}^N (x - \mu^*)(x - \mu^*)^T = \begin{bmatrix} 0.04 & 0 \\ 0 & 0.03 \end{bmatrix}$$

# GMM (Gaussian Mixture Model)

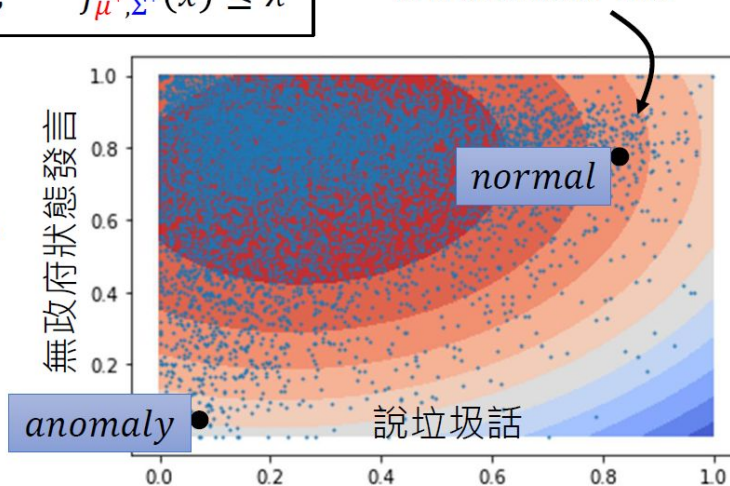
$$f_{\mu^*, \Sigma^*}(x) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma^*|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu^*)^T \Sigma^{*-1} (x - \mu^*) \right\}$$

$$\mu^* = \begin{bmatrix} 0.29 \\ 0.73 \end{bmatrix} \quad \Sigma^* = \begin{bmatrix} 0.04 & 0 \\ 0 & 0.03 \end{bmatrix}$$

$$f(x) = \begin{cases} \text{normal}, & f_{\mu^*, \Sigma^*}(x) > \lambda \\ \text{anomaly}, & f_{\mu^*, \Sigma^*}(x) \leq \lambda \end{cases}$$

$\lambda$  is a contour line

The colors represents  
the value of  $f_{\mu^*, \Sigma^*}(x)$



# GMM (Gaussian Mixture Model)

```
1  from keras.datasets import mnist
2  import cv2
3  import numpy as np
4
5  (x_train, y_train), (x_test, y_test) = mnist.load_data()
6  x_ok = x_train[y_train == 1] # 6742 筆
7  x_test = x_test[(y_test == 7) | (y_test == 1)] # 1135 筆 "1", 1028 筆 "7"
8  y_test = y_test[(y_test == 7) | (y_test == 1)]
9
10 def reshape_x(x):
11     new_x = np.empty((len(x), 56, 56))
12     for i, e in enumerate(x):
13         new_x[i] = cv2.resize(e, (56, 56))
14
15     new_x = np.expand_dims(new_x, axis=-1)
16     new_x = np.repeat(new_x, 3, axis=-1)
17     return new_x
18
19 x_ok = reshape_x(x_ok)
20 x_test = reshape_x(x_test)
```

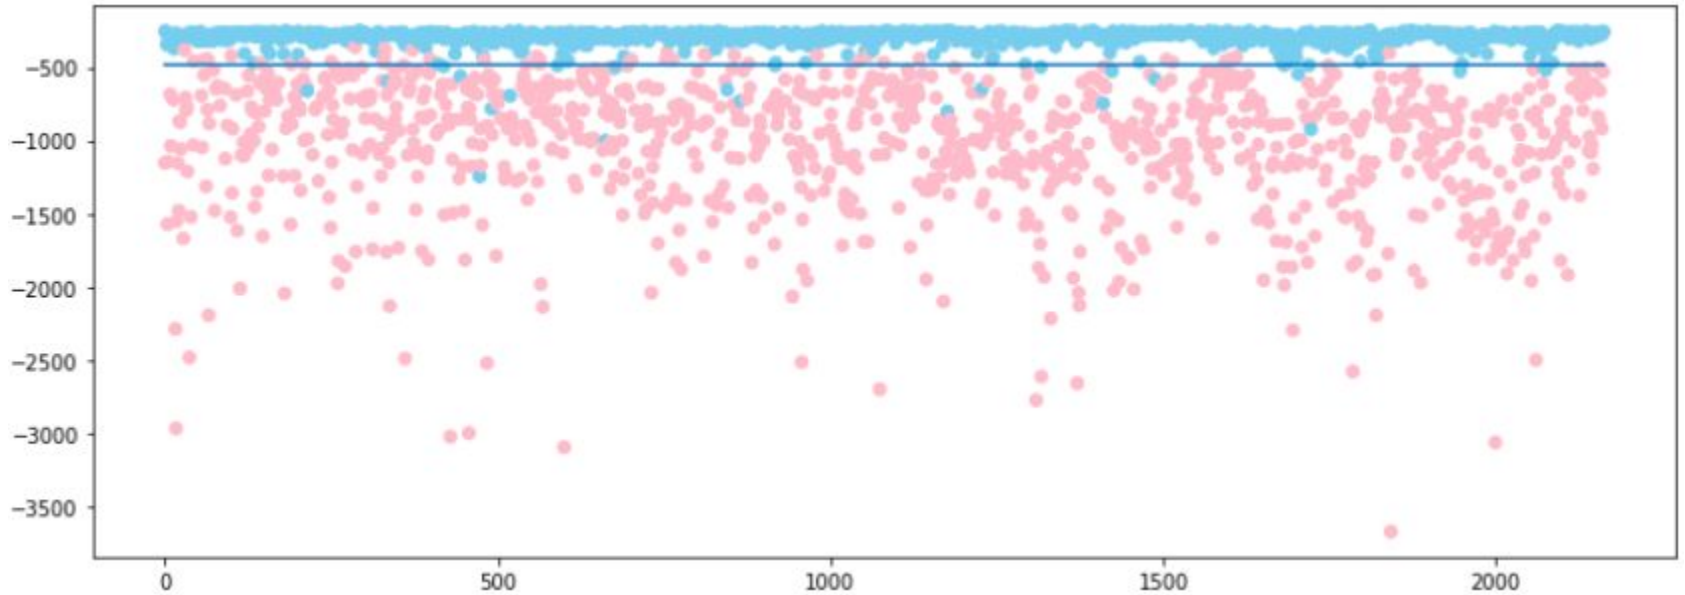
# GMM (Gaussian Mixture Model)

```
features = model.predict(x_ok)
gmm.fit(features)
OKscore = gmm.score_samples(features)
thred = OKscore.mean() - 3 * OKscore.std()

test_features = model.predict(x_test)
score = gmm.score_samples(test_features)

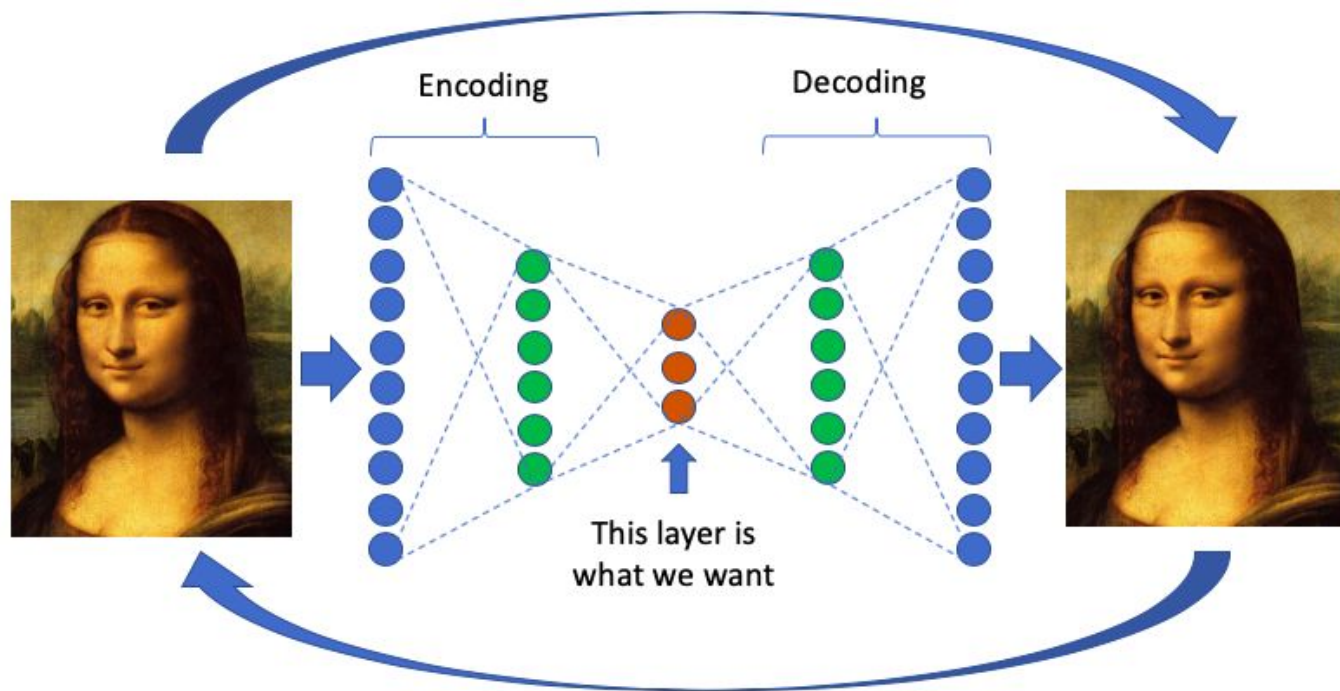
print('normal accuracy: %.2f' % (len(score[(y_test == 1) & (score > thred)]) / 1135))
print('abnormal accuracy: %.2f' % (len(score[(y_test == 7) & (score < thred)]) / 1028))
      normal accuracy: 0.98
      abnormal accuracy: 0.96
```

# GMM (Gaussian Mixture Model)





# Auto-Encoder





# Auto-Encoder

Examples:

<https://www.kaggle.com/tikedameu/anomaly-detection-with-autoencoder-pytorch>

<https://towardsdatascience.com/anomaly-detection-with-autoencoder-b4cdce4866a6>


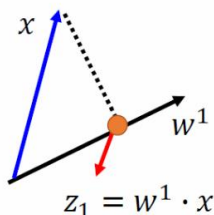
# PCA

## PCA

$$z = Wx$$

Reduce to 1-D:

$$z_1 = w^1 \cdot x$$



Project all the data points  $x$  onto  $w^1$ , and obtain a set of  $z_1$

We want the variance of  $z_1$  as large as possible

$$\text{Var}(z_1) = \sum_{z_1} (z_1 - \bar{z}_1)^2 \quad \|w^1\|_2 = 1$$

## PCA

$$z = Wx$$

Reduce to 1-D:

$$z_1 = w^1 \cdot x$$

$$z_2 = w^2 \cdot x$$

$$W = \begin{bmatrix} (w^1)^T \\ (w^2)^T \\ \vdots \end{bmatrix}$$

Orthogonal matrix

Project all the data points  $x$  onto  $w^1$ , and obtain a set of  $z_1$

We want the variance of  $z_1$  as large as possible

$$\text{Var}(z_1) = \sum_{z_1} (z_1 - \bar{z}_1)^2 \quad \|w^1\|_2 = 1$$

We want the variance of  $z_2$  as large as possible

$$\text{Var}(z_2) = \sum_{z_2} (z_2 - \bar{z}_2)^2 \quad \|w^2\|_2 = 1$$

$$w^1 \cdot w^2 = 0$$

# PCA

```
# 30 principal components
from sklearn.decomposition import PCA

n_components = 30
whiten = False
random_state = 2018


pca = PCA(n_components=n_components, whiten=whiten, \
          random_state=random_state)

X_train_PCA = pca.fit_transform(X_train)
X_train_PCA = pd.DataFrame(data=X_train_PCA, index=X_train.index)

X_train_PCA_inverse = pca.inverse_transform(X_train_PCA)
X_train_PCA_inverse = pd.DataFrame(data=X_train_PCA_inverse, \
                                   index=X_train.index)

scatterPlot(X_train_PCA, y_train, "PCA")
```

# PCA

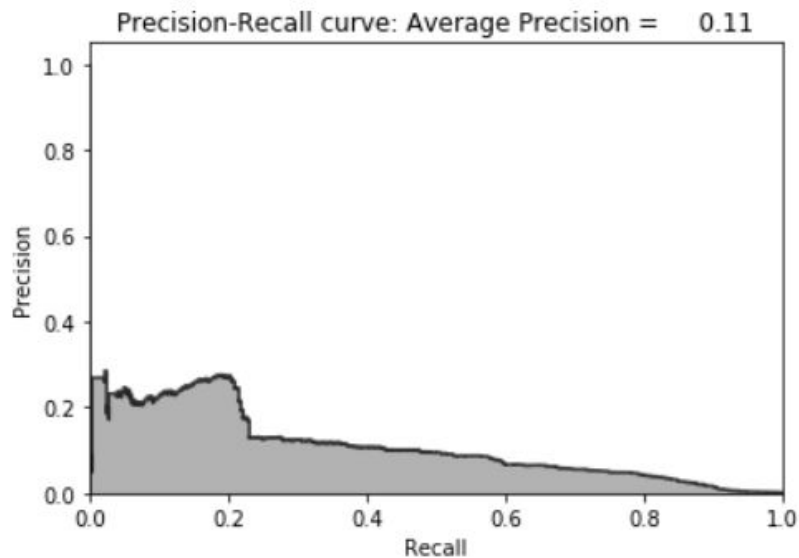


# PCA

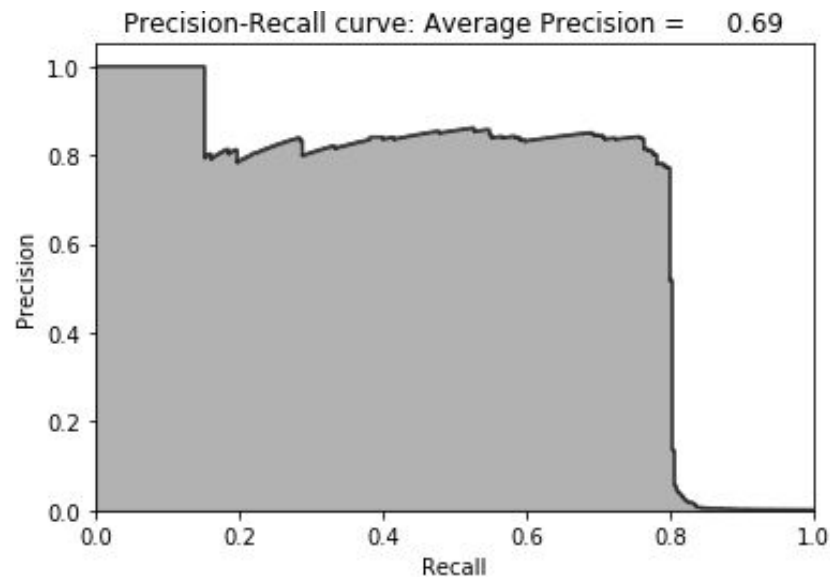
```
def anomalyScores(originalDF, reducedDF):  
    loss = np.sum((np.array(originalDF) - np.array(reducedDF))**2, axis=1)  
    loss = pd.Series(data=loss, index=originalDF.index)  
    loss = (loss - np.min(loss)) / (np.max(loss) - np.min(loss))  
    return loss
```

```
anomalyScoresPCA = anomalyScores(X_train, X_train_PCA_inverse)  
preds = plotResults(y_train, anomalyScoresPCA, True)
```

# PCA



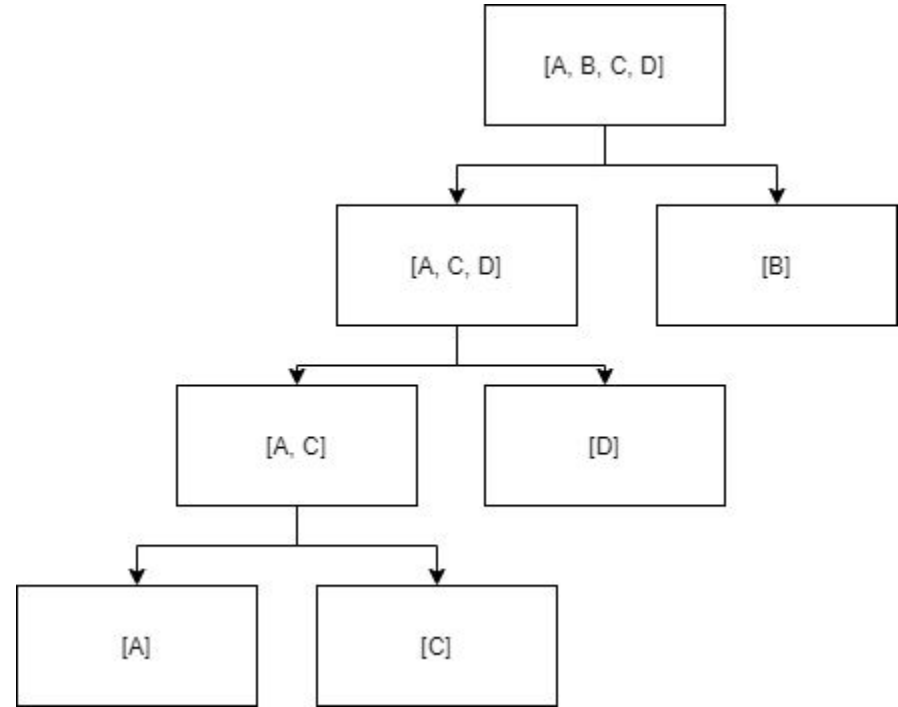
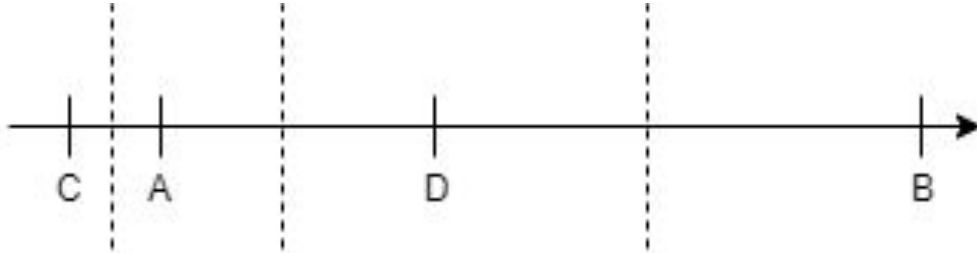
n\_components = 30



n\_components = 27

# Isolation Forest

- Use tree like structure to split data



# Isolation Forest

```
# importing libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from pylab import savefig
from sklearn.ensemble import IsolationForest

# Generating data ----

rng = np.random.RandomState(42)

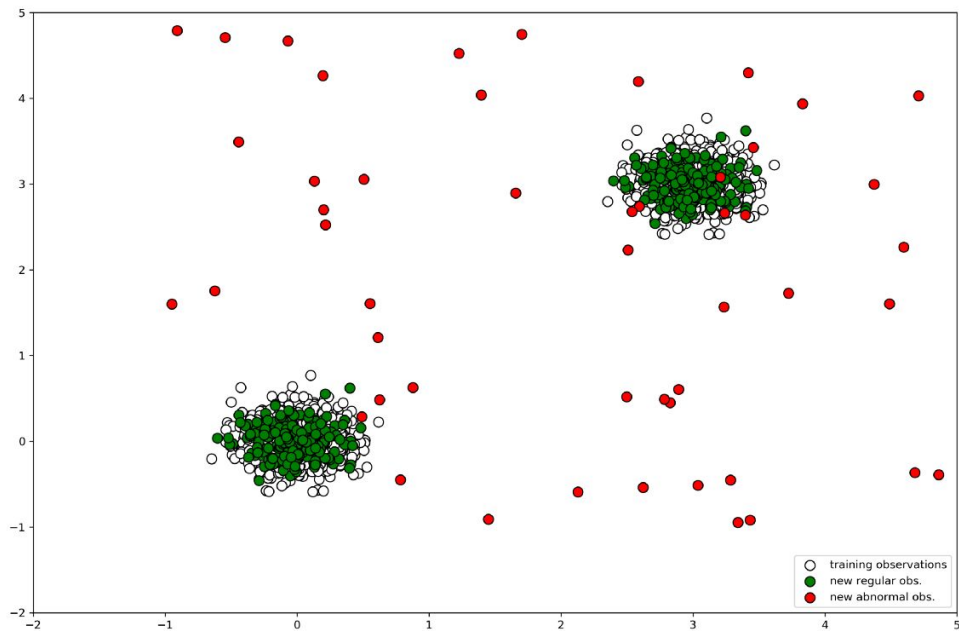
# Generating training data
X_train = 0.2 * rng.randn(1000, 2)
X_train = np.r_[X_train + 3, X_train]
X_train = pd.DataFrame(X_train, columns = ['x1', 'x2'])

# Generating new, 'normal' observation
X_test = 0.2 * rng.randn(200, 2)
X_test = np.r_[X_test + 3, X_test]
X_test = pd.DataFrame(X_test, columns = ['x1', 'x2'])

# Generating outliers
X_outliers = rng.uniform(low=-1, high=5, size=(50, 2))
X_outliers = pd.DataFrame(X_outliers, columns = ['x1', 'x2'])
```



# Isolation Forest



# Isolation Forest

```
# Isolation Forest ----  
  
# training the model  
clf = IsolationForest(max_samples=100, random_state=rng)  
clf.fit(X_train)  
  
# predictions  
y_pred_train = clf.predict(X_train)  
y_pred_test = clf.predict(X_test)  
y_pred_outliers = clf.predict(X_outliers)
```

# Isolation Forest

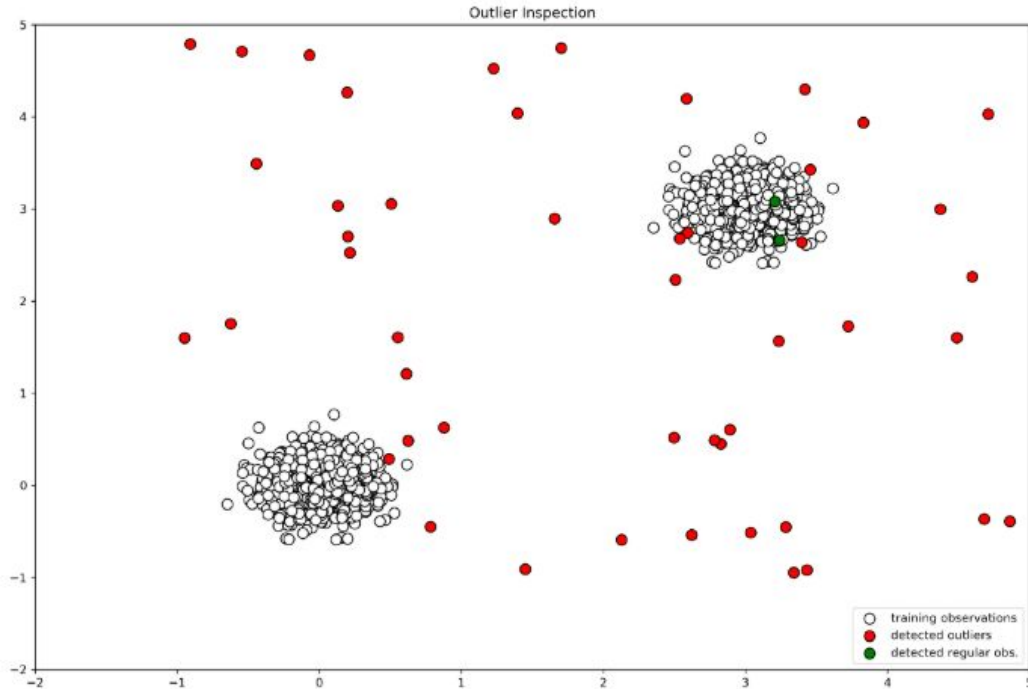


Figure 3 Inspecting outlier classification

<https://towardsdatascience.com/outlier-detection-with-isolation-forest-3d190448d45e>

# Summary of classic method

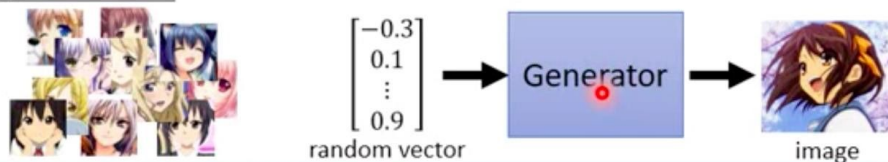
- With Classifier
- GMM (Gaussian Mixture Model)
- Auto-Encoder
- PCA
- Isolation Forest

Anomaly detection on image

# Typical GANs

## Three Categories of GAN

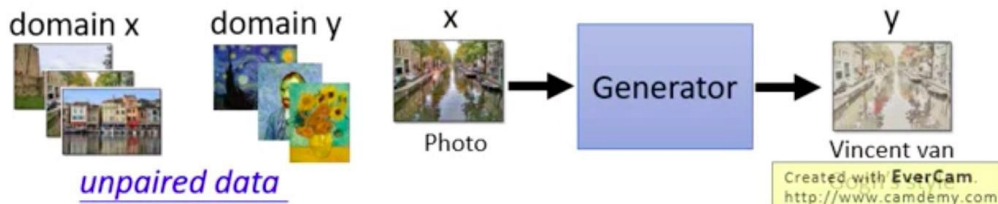
### 1. Typical GAN



### 2. Conditional GAN



### 3. Unsupervised Conditional GAN



# Typical GANs

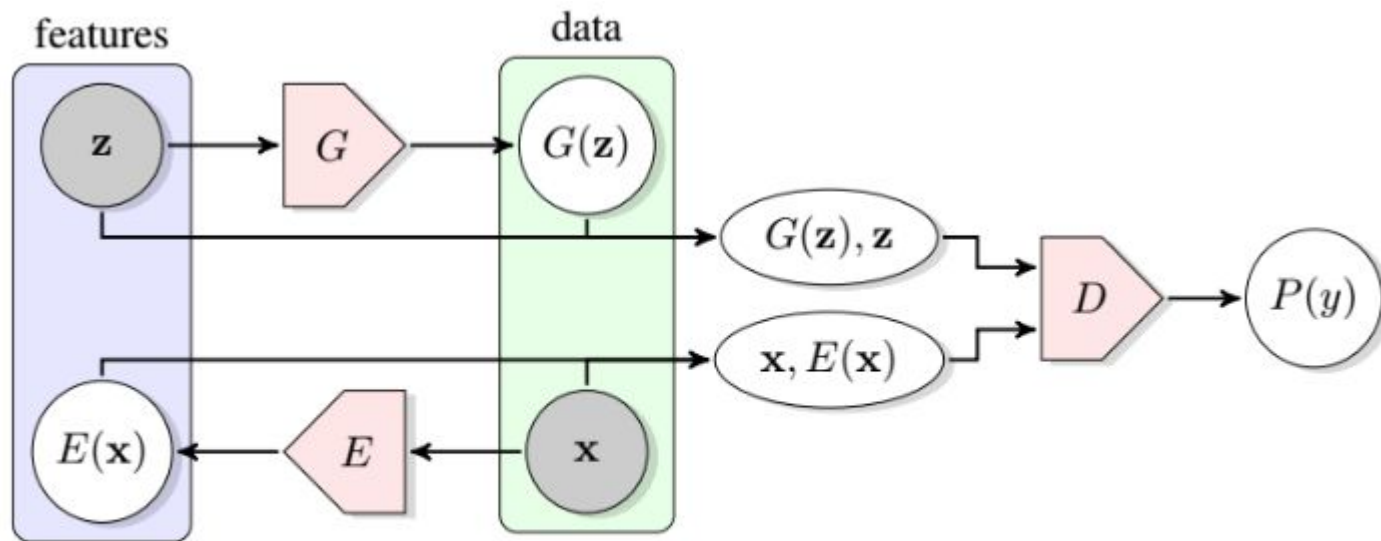
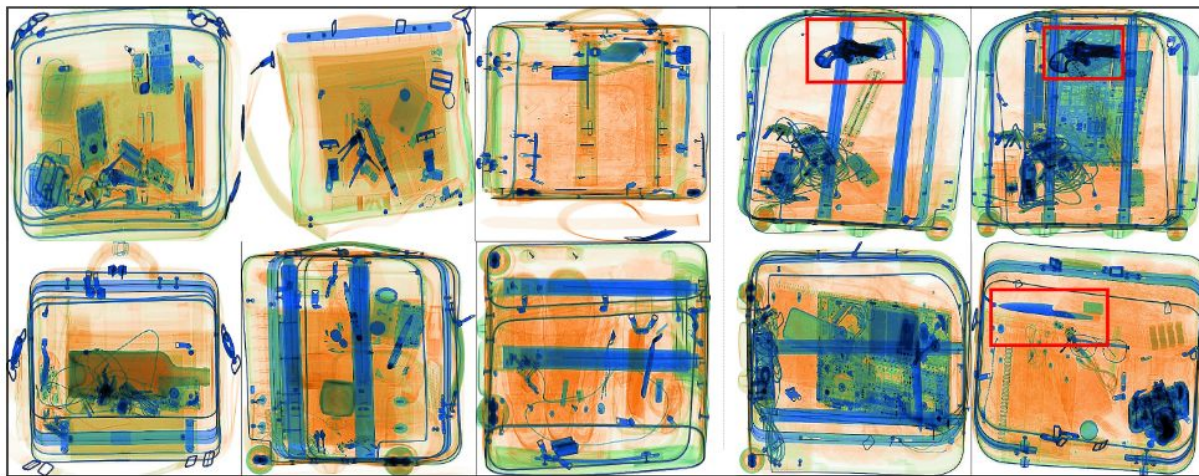


Figure 1. The structure of BiGAN proposed in (Donahue et al., 2016).

# Example



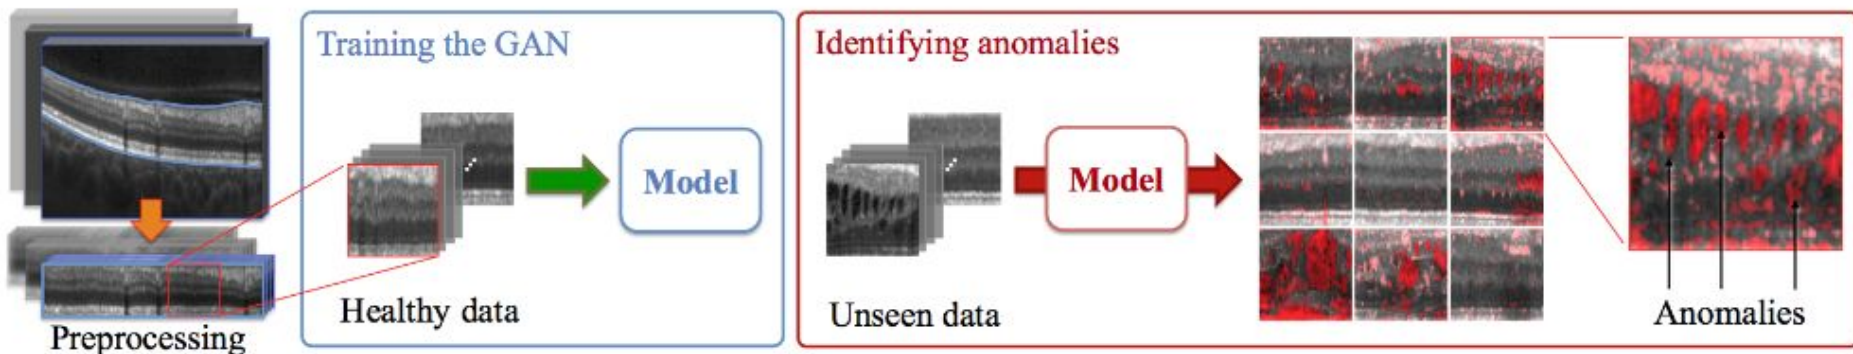
(a) Normal Data (X-ray Scans)

(b) Normal + Abnormal Data (X-ray Scans)



# AnoGAN

- Train a standard GAN only on positive samples



# AnoGAN

- Anomaly Score

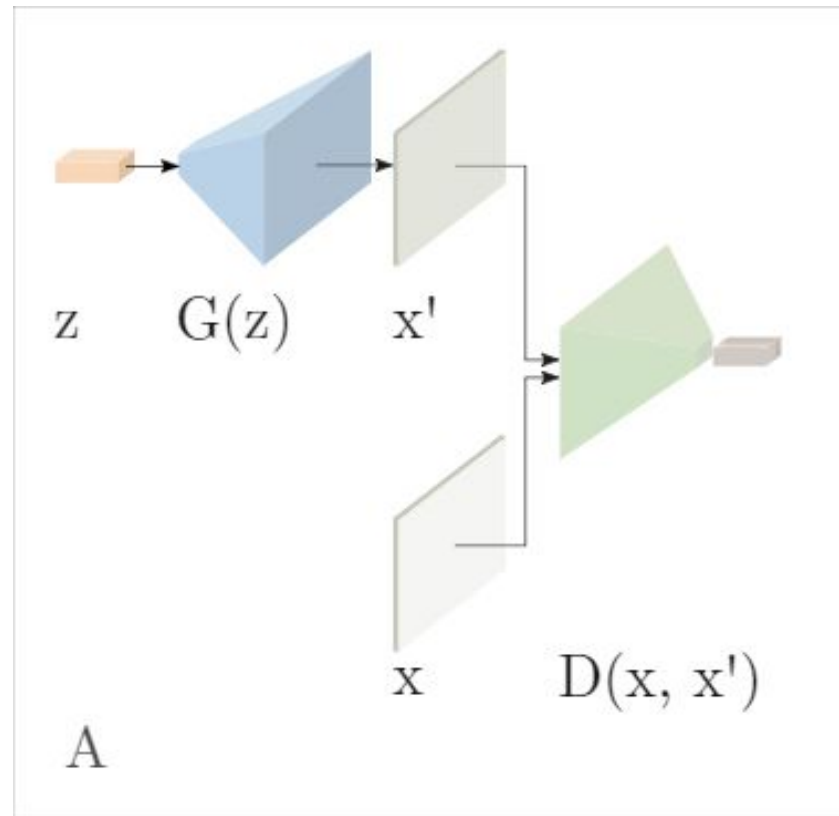
$$A(\mathbf{x}) = \mathcal{L}(\mathbf{z}_\Gamma)$$

- for  $\gamma = 1, 2, \dots, \Gamma$  find proper  $\mathbf{z}$

$$\mathcal{L}_R(\mathbf{z}_\gamma) = \|\mathbf{x} - G(\mathbf{z}_\gamma)\|_1$$

$$\mathcal{L}_D(\mathbf{z}_\gamma) = \|\mathbf{f}(\mathbf{x}) - \mathbf{f}(G(\mathbf{z}_\gamma))\|_1$$

$$\mathcal{L}(\mathbf{z}_\gamma) = (1 - \lambda) \cdot \mathcal{L}_R(\mathbf{z}_\gamma) + \gamma \cdot \mathcal{L}_D(\mathbf{z}_\gamma)$$



# AnoGAN

- Pros

- Showed that GANs can be used for anomaly detection
- Introduced a new mapping scheme from latent space to input data space.
- Used the same mapping scheme to define an anomaly score.

- Cons

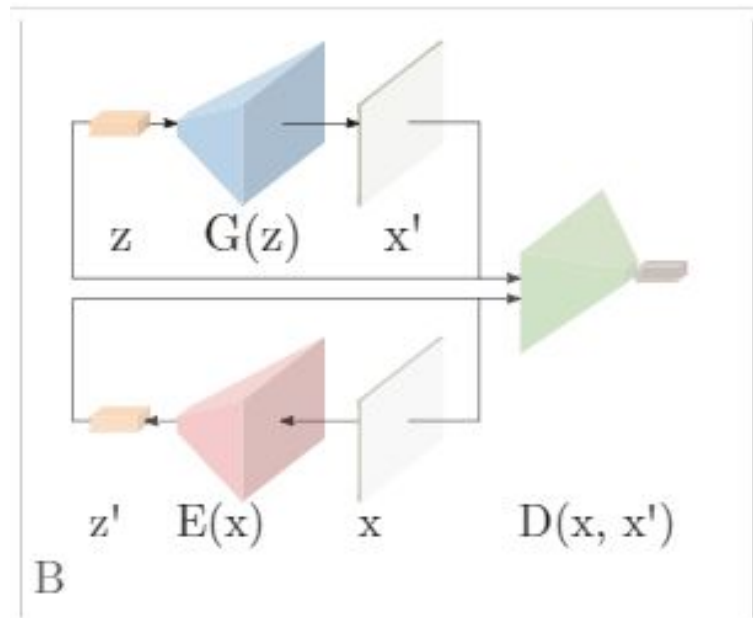
- Requires  $\Gamma$  optimization steps for every new input: bad test-time performance
- The GAN objective has not been modified to take into account the need for the inverse mapping learning.
- The anomaly score is difficult to interpret, not being in the probability range.

# EGBAD (Efficient GAN-Based Anomaly Detection)

- Train a Bi-GAN only on positive samples

$$\min_{G,E} \max_D V(D, G, E) =$$

$$\mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\mathbb{E}_{\mathbf{z} \sim p_{E(\mathbf{z}|\mathbf{x})}} [\log D(\mathbf{x}, \mathbf{z})]] +$$
$$\mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\mathbb{E}_{\mathbf{x} \sim p_G(\mathbf{x}|\mathbf{z})} [\log(1 - D(\mathbf{x}, \mathbf{z}))]].$$



# EGBAD (Efficient GAN-Based Anomaly Detection)

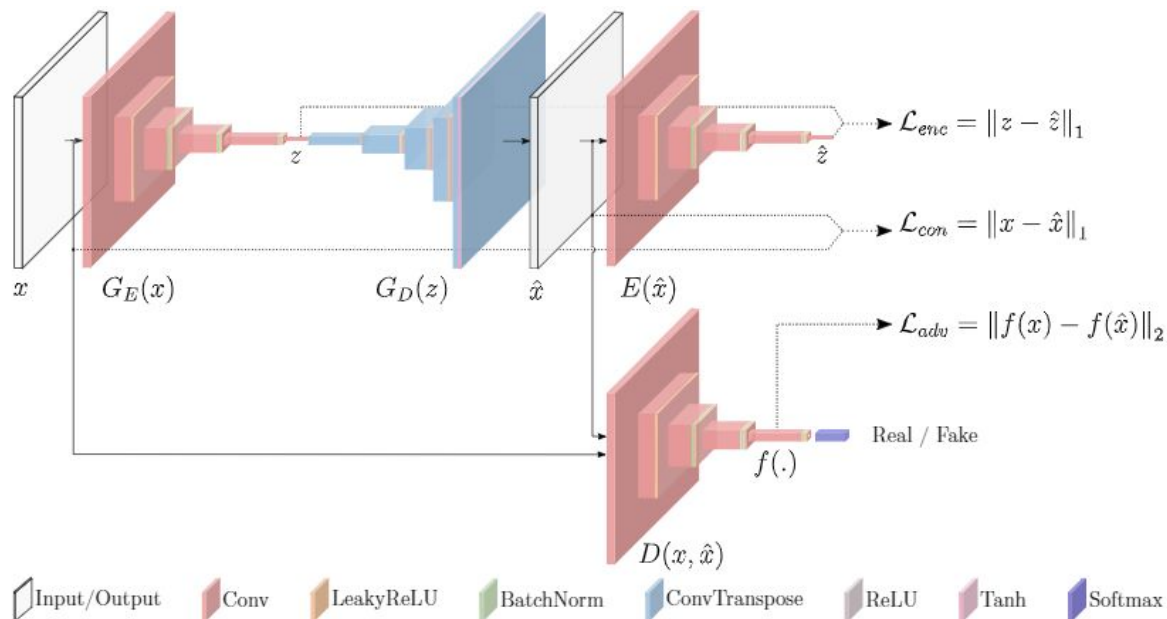
- Pros
  - The Encoder E can learn how to encode image while adversarial training.
  - Therefore, it can bypass  $\Gamma$  optimization steps of AnoGAN while calculating anomaly score.

# GANomaly

- Generator is composed of encoder  $G_E$ , decoder  $G_D$ , and encoder  $E$
- Trained on only normal data.

- Anomaly score

$$\mathcal{A}(\mathbf{x}) = \|G_E(\mathbf{x}) - E(G(\mathbf{x}))\|_2$$



# GANomaly

- Pros

- An encoder is learned during the training process, so it can bypass the  $\Gamma$  optimization.
- Using an autoencoder like architecture (no use of noise prior) makes the entire learning process faster.
- The contextual loss can be used to localize the anomaly.

- Cons

- Defines a new anomaly score.
- It allows to detect anomalies both in the image space and in the latent space, but the results couldn't match:
  - a higher anomaly score, that's computed only in the latent space, can be associated with a generated sample with a low contextual loss value and thus very similar to the input - and vice versa.


# Evaluation Metric

- TPR (True Positive Rate)

$$\frac{TP}{TP + FN}$$

- FPR (False Positive Rate)


$$\frac{FP}{TN + FP}$$






# Comparison

**BiGAN/EGBAD and GANomaly performance comparison on MNIST and Fashion-MNIST datasets**


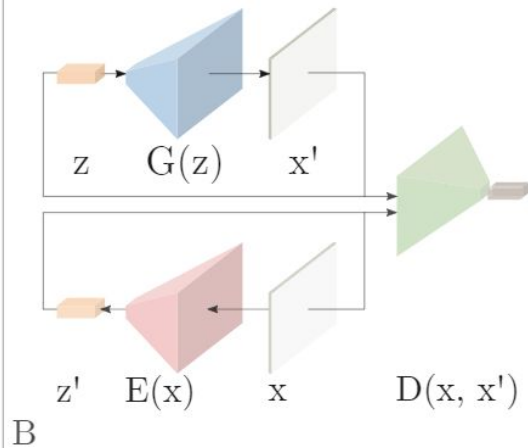
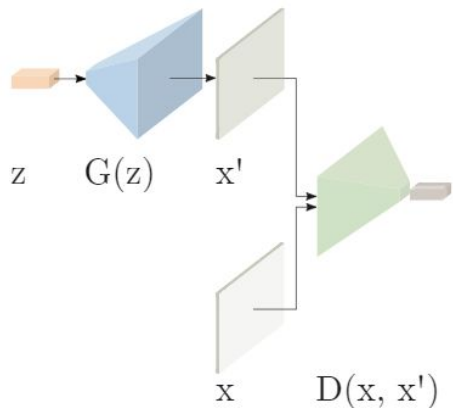


(a) MNIST anomalous classes



(b) Fashion-MNIST anomalous classes

# Summary of GANs



AnoGAN

EGBAD

GANomaly

# Reference

<https://arxiv.org/pdf/1711.09325.pdf>

<https://arxiv.org/pdf/1809.10816.pdf>


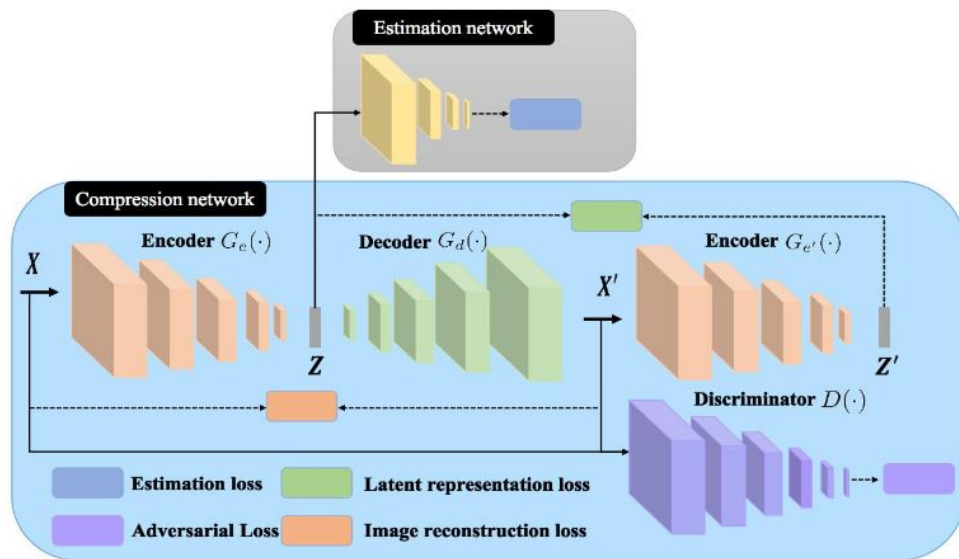
<https://arxiv.org/pdf/1812.02288.pdf>

<https://arxiv.org/pdf/1901.08954.pdf>

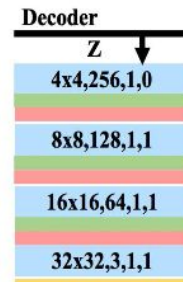
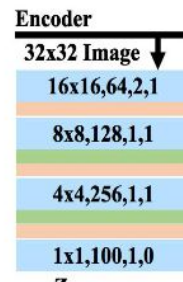
<https://arxiv.org/pdf/1905.13147.pdf>

# Anomaly detection on Audio

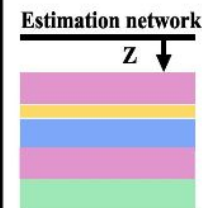
# GMGAN (Gaussian Mixture GAN)



Fake / Real



32x32 Image



# GMGAN (Gaussian Mixture GAN)

- Adversarial loss

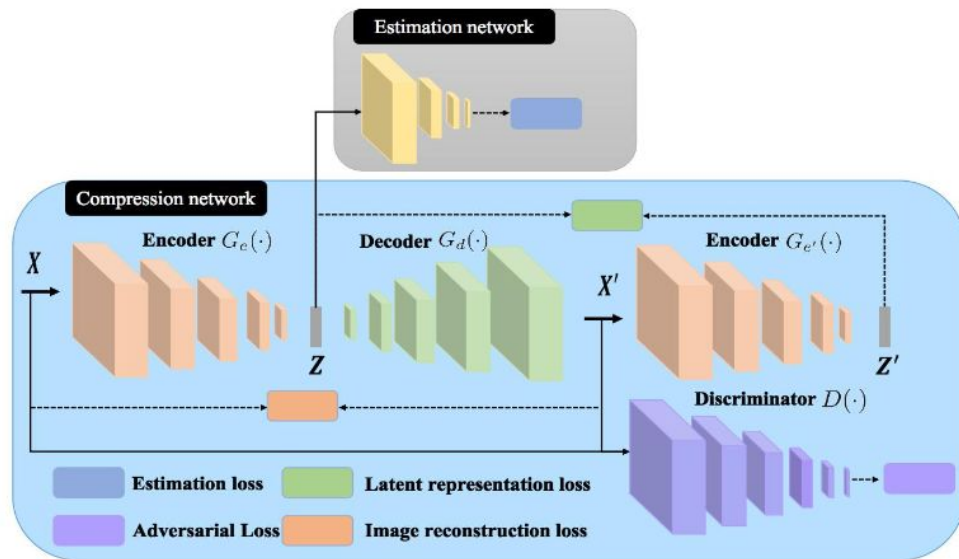
$$\mathcal{L}_{adv} = \min_G \max_D (E_{\mathbf{x} \sim p_{\mathbf{x}}} [\log(D(\mathbf{x}))] \\ + E_{\mathbf{x} \sim p_{\mathbf{x}}} [\log(1 - D(G(\mathbf{x})))]).$$

- Image reconstruction loss

$$\mathcal{L}_{irec} = \mathbb{E}_{x \sim p_{\mathbf{x}}} \|x - G(x)\|_1$$

- Latent representation loss

$$\mathcal{L}_{zrec} = \mathbb{E}_{x \sim p_{\mathbf{X}}} \|G_e(x) - G_{e'}(x')\|_2$$

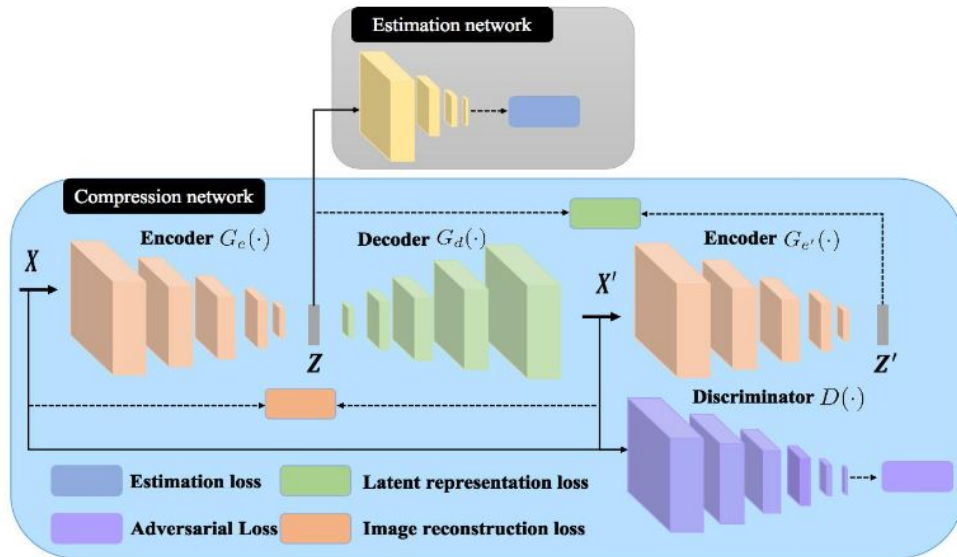


# GMGAN (Gaussian Mixture GAN)

- Estimation loss
- $\hat{\gamma} = \text{softmax}(MLN(\mathbf{z}; \theta_m))$
- $\hat{\gamma}$  is a K-dimensional vector and  $\hat{\gamma}_k$  denotes the input belongs to the  $k^{\text{th}}$  distribution.
- Calculate the component of  $k^{\text{th}}$  mixture.

$$\hat{\alpha}_k = \frac{1}{n} \sum_{i=1}^n \hat{\gamma}_{ik}; \hat{u}_k = \frac{\sum_{i=1}^n \hat{\gamma}_{ik} z_i}{\sum_{i=1}^n \hat{\gamma}_{ik}},$$

$$\hat{\Sigma}_k = \frac{\sum_{i=1}^n \hat{\gamma}_{ik} (z_i - \hat{u}_k)(z_i - \hat{u}_k)^T}{\sum_{i=1}^n \hat{\gamma}_{ik}},$$

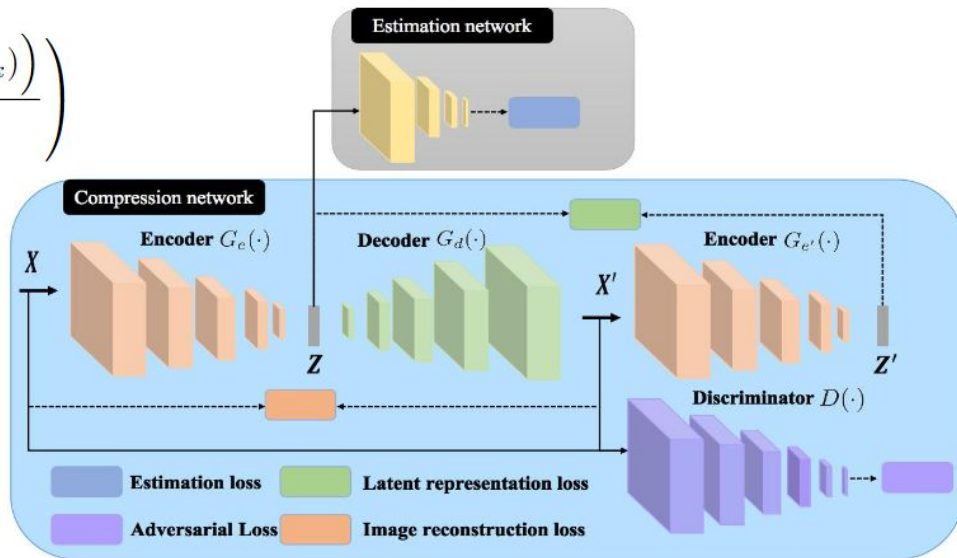


# GMGAN (Gaussian Mixture GAN)

- Estimation loss


$$E(\mathbf{z}) = -\log \left( \sum_{k=1}^K \hat{\alpha}_k \frac{\exp \left( -\frac{1}{2} (\mathbf{z} - \hat{\mu}_k)^T \hat{\Sigma}_k^{-1} (\mathbf{z} - \hat{\mu}_k) \right)}{\sqrt{|2\pi \Sigma_k|}} \right)$$

$$\mathcal{L}_{es} = \lambda_1 \sum_{i=1}^N E(z_i) + \lambda_2 \sum_{k=1}^K \sum_{j=1}^d \frac{1}{\hat{\Sigma}_{jj}^k},$$





# GMGAN (Gaussian Mixture GAN)



**Fig. 2.** Spectrogram of one second sequence containing office background audio and glass break audio.

# GMGAN (Gaussian Mixture GAN)

Scene	CAE [9]	WaveNet [19]	Proposed method
beach	0.69	0.72	<b>0.80</b>
bus	0.79	0.83	<b>0.89</b>
cafe/restaurant	0.69	<b>0.76</b>	<b>0.76</b>
car	0.79	0.82	<b>0.93</b>
city center	0.75	0.82	<b>0.83</b>
forest path	0.65	0.72	<b>0.77</b>
grocery store	0.71	0.77	<b>0.83</b>
home	<b>0.69</b>	<b>0.69</b>	<b>0.69</b>
library	0.59	0.67	<b>0.85</b>
metro station	0.74	0.79	<b>0.81</b>
office	0.78	0.78	<b>0.80</b>
park	0.70	0.80	<b>0.89</b>
residential area	0.73	<b>0.78</b>	<b>0.78</b>
train	0.82	0.84	<b>0.92</b>
tram	0.80	0.87	<b>0.94</b>

# Outline

- What is Anomaly Detection
- Classic Method
  - With Classifier
  - GMM (Gaussian Mixture Model)
  - Auto-Encoder
  - PCA
  - Isolation Forest
  - Summary
- Anomaly Detection with GAN
  - AnoGAN
  - EGBAD
  - GANomaly
  - Summary
- Anomaly Detection on Audio
  - GMGAN

# More Reference

Anomaly detection on self-driving

<http://taoxie.cs.illinois.edu/publications/tii17-safedrive.pdf>

<https://arxiv.org/pdf/2004.09496.pdf>

<https://arxiv.org/pdf/2004.12581.pdf>